

**GOEXPERIMENT=arena**

# Vale a pena usar arenas?

# O que eu venho fazendo?



stone



MANNING

Packt

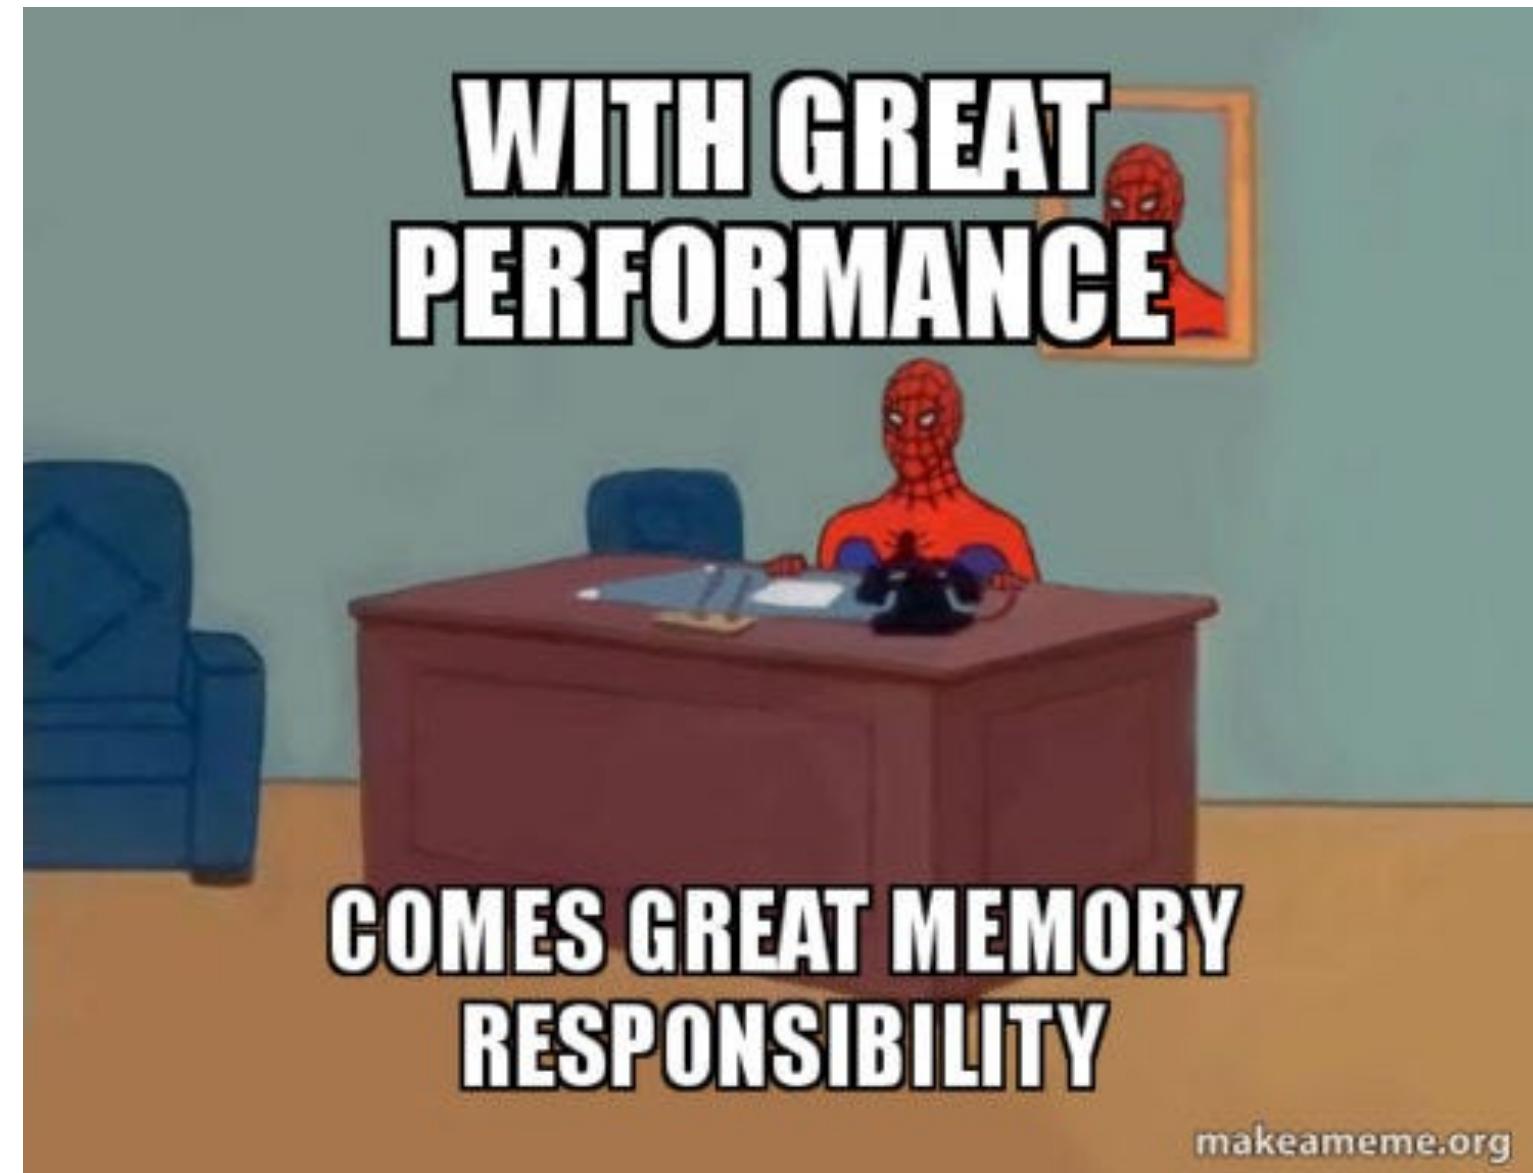


Me, who forgot a  
free() in a loop

Malloc

Memory: 104390 MB

# Memória não é infinita



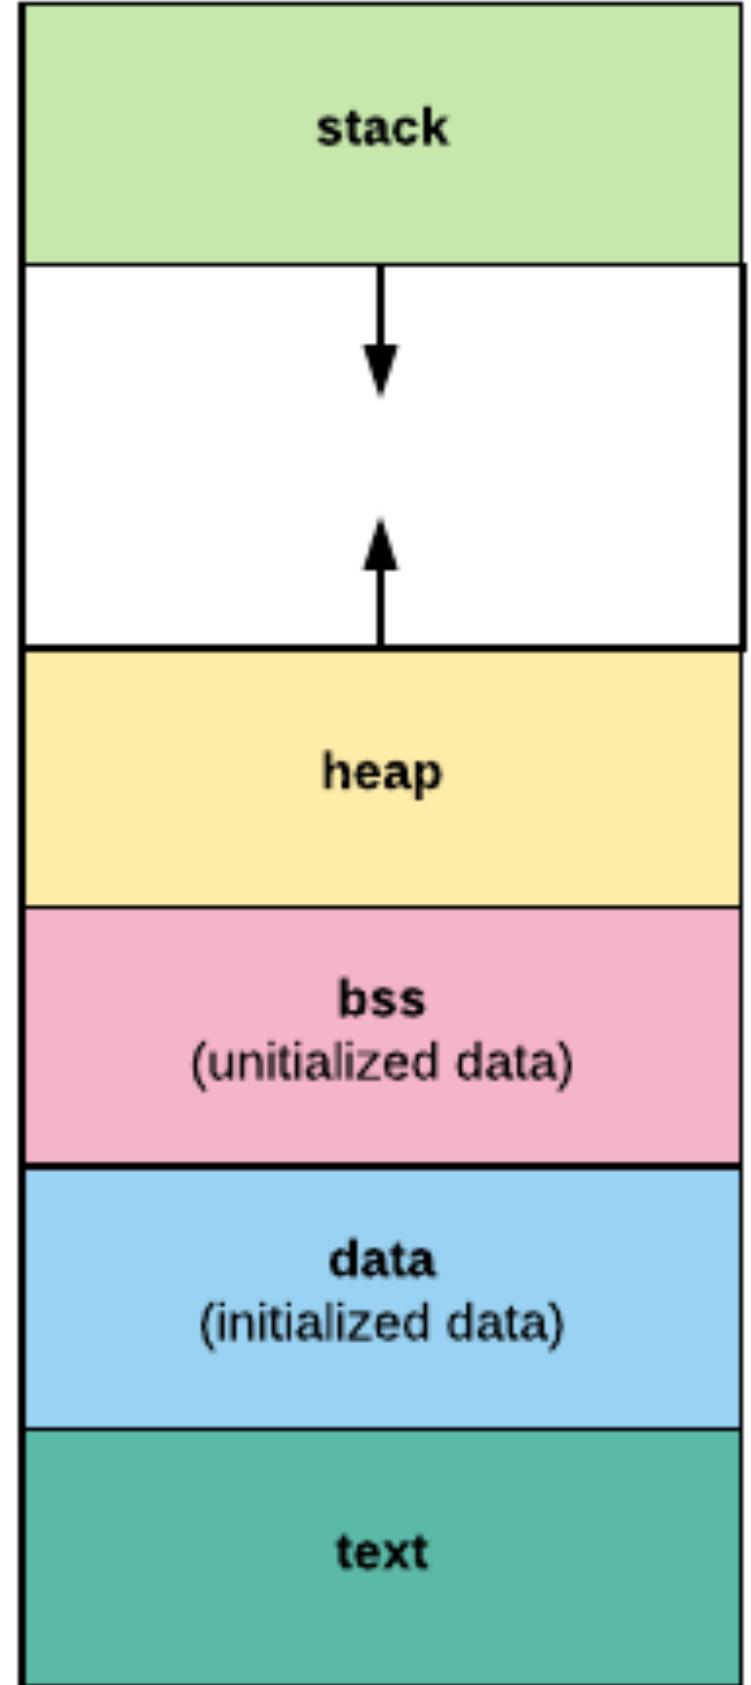
# AGENDA

- 1 - Memory Layout
- 2 - Garbage Collector
- 3 - Go Memory Arena
- 4 - Case Study
- 5 - Guidelines

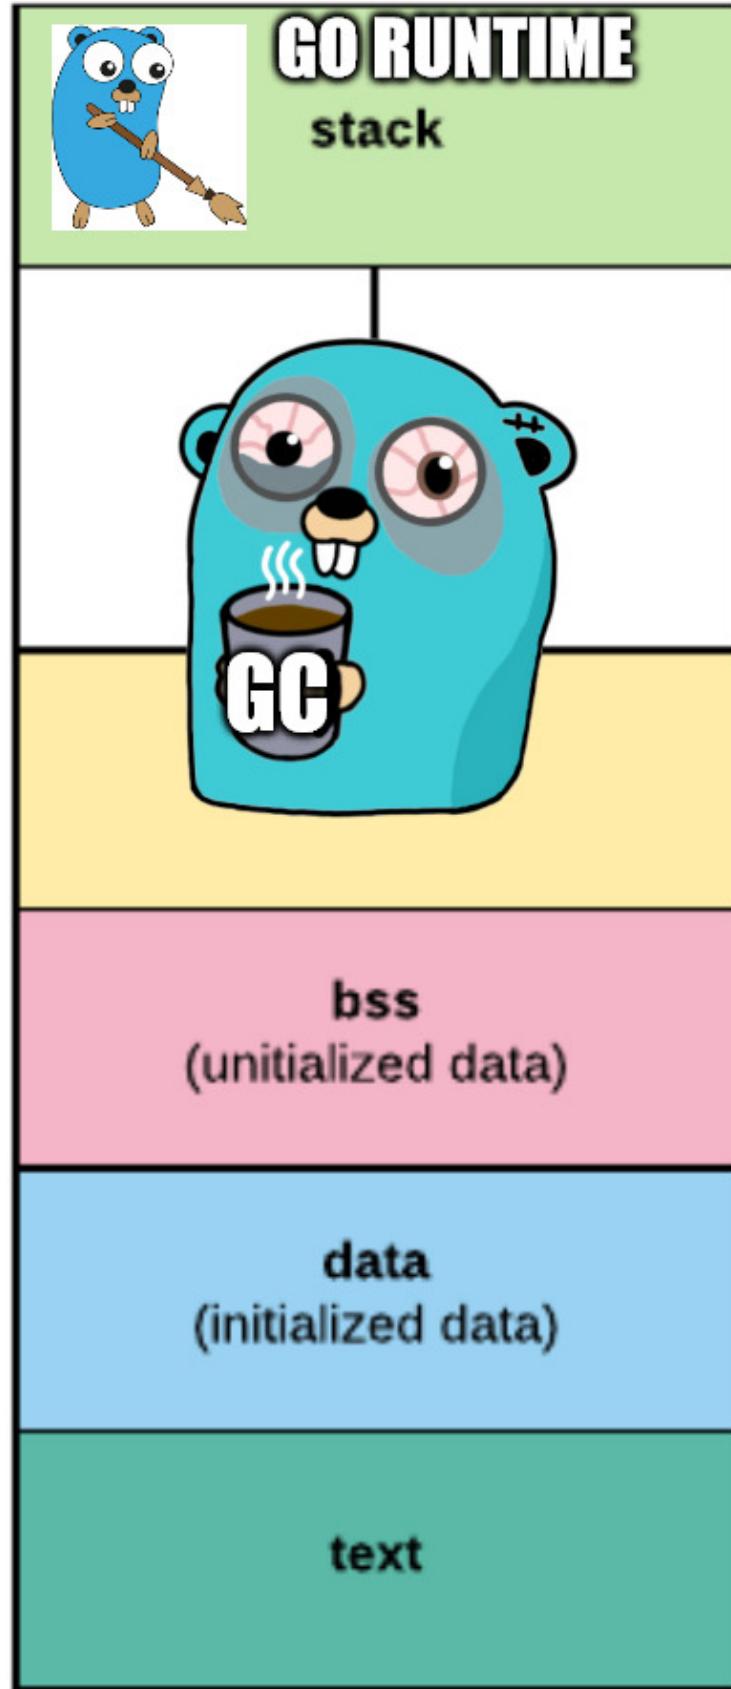
# 1 - Memory Layout

# Memory Layout

- » Known size objects -> Runtime managed
- » Unknown size objects -> GC Managed



# GC vs. Runtime



# 2 - Garbage Collector

Maior foco no programa e menos na  
memória

# Evitar accidentes



# Responsabilidades

- » Rastrear alocações na heap
- » Liberar alocações que não são mais necessárias
- » Manter as alocações em uso

# Memory inference

## O que liberar?

# Memory inference

- » Tracing GC
- » Reference Counting GC

# Tracing GC

**Tracing Garbage Collection** is a form of automatic memory management that consists of determining which objects should be deallocated ("garbage collected") by tracing which objects are *reachable* by a chain of references from certain "root" objects, and considering the rest as "garbage" and collecting them.

# The Journey of Go's Garbage Collector

<https://go.dev/blog/ismmkeynote>

# Timeline

- » 1.0 - Mark and Sweep (300ms)
- » 1.5 - Tricolor, Concurrent, Mark and Sweep (10ms)
- » latest - Optimizations (0,5ms)

# GC Phases

- » Mark setup - STW
- » Marking - Concurrent
- » Mark Termination - STW

# GC Phases - Mark

## Setup

- » Liga Write Barrier
- » Todas as goroutines devem estar paradas.



# GC Phases - Marking

- » Procura pelos ponteiros nas roots para a heap
- » Percorre o grafo da heap desses ponteiros
- » Marca os valores na heap como em uso

# GC Phases - Concurrent Marking

Concurrent:

- 25% of the available CPU for itself
- Até +5% com mark assist

# GC Phases – Mark Assist



# GC Phases - Mark Termination

- » Turn off the Write Barrier
- » Clean up tasks
- » Next collection goal is calculated



# Sweep

Freeing Heap Memory



# Fine tuning GOGC (env var)



# Default GOGC=100

Próxima execução quando o tamanho da heap aumentar 100%

# Heap Goal

Se a coleta terminar com 10MB na heap a próxima irá executar quando a heap chegar a 10MB de novas alocações.

# GC Pacer

Quando começar uma nova coleta?

# Collection Cadence

Estima quanto tempo irá levar para finalizar a coleta.

**+ GC = + Latency** 

- » Rouba capacidade da CPU (no mínimo 25%)
- » STW Latency (0% for work)

# Resumo do estado atual da GC

- » 25% CPU durante a coleta
- » Heap 2x live heap (ou max heap)
- » 2 STW com tempos menores que 0,5ms por ciclo de coleta
- » Em steady state temos poucos Mark Assists

# Ajude o GC a te ajudar

P1 - Identifique e remova alocações desnecessárias

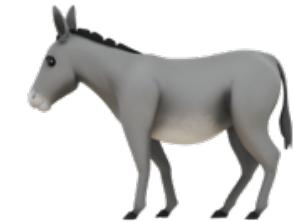
P2:

- Mantenha a heap tão pequena quanto possível
- Encontre o pace ideal
- Fique dentro dos objetivos para cada coleta
- Minimize o tempo da coleta, STW e Mark Assists

# Problema

Alocar um grande quantidade de objetos em uma operação em um período de tempo bem definido.

# GOGC=off ?



OOM 

# Soft limit?

## GOMEMLIMIT (env var)

# Latency vs. memory footprint



imgflip.com

JAKE-CLARK.TUMBLR

# No manual management



# Memory Region

- » AED Free Storage Package (1967) - zones
- » PL/I (1976) - AREA data type
- » C (1990) - Arenas
- » Apache HTTP Server - Pools
- » Postgres - Memory Context

*[https://en.wikipedia.org/wiki/Region-basedmemorymanagement](https://en.wikipedia.org/wiki/Region-based_memory_management)*

# 3 - Go Memory Arena

# proposal: arena: new package providing memory arenas

<https://github.com/golang/go/issues/51317>

*This proposal is on hold indefinitely due to serious API concerns. The GOEXPERIMENT=arena code may be changed incompatibly or removed at any time, and we do not recommend its use in production.*

# Objetivos

- » Alocar de uma vez uma região de memória (chunk) para nossos objetos.
- » Poder desalocar todo essa região de uma vez só.
- » Não introduzir os problemas da gestão manual.

# Como usar?

GOEXPERIMENT=arena

```
import "arena"
```

# Arena - crash course

- » Criar uma arena
- » Fazer operações usando arena para alocar memória onde for necessário.
- » Liberar a arena.

# Criando uma nova Arena

```
func NewArena() Arena
```

---

```
mem := arena.NewArena()
```

# Incluindo objetos na Arena

```
func New[T any](a *Arena) *T
```

---

```
mem := arena.NewArena()  
p := arena.New[Person](mem)
```

# Colocando objetos na heap

```
func Clone[T any](s T) T
```

---

```
mem := arena.NewArena()  
p1 := arena.New[Person](mem) // arena-allocated  
p2 := arena.Clone(p1) // heap-allocated
```

# Liberando a memória

**func (a \*Arena) Free()**

-----

**mem.Free()**

# Criando slices

```
func MakeSlice[T any](a *Arena, len, cap int) []T
```

-----

```
mem := arena.NewArena()
```

```
slice := arena.MakeSlice[string](mem, 100, 100)
```

# Address sanitizer

```
type T struct {
    Num int
}

func main() {
    mem := arena.NewArena()
    o := arena.New[T](mem)
    mem.Free()
    o.Num = 123 // 🤡
}
```

# Address sanitizer (Asan)

go run -asan main.go

accessed data from freed user arena 0x40c0007ff7f8

# Diferença na abordagem

Alocando na heap ->  $O(N)$  deallocs

Alocando na arena ->  $O(1)$  dealloc

# Performance?

Performance não se adivinha, se mede.

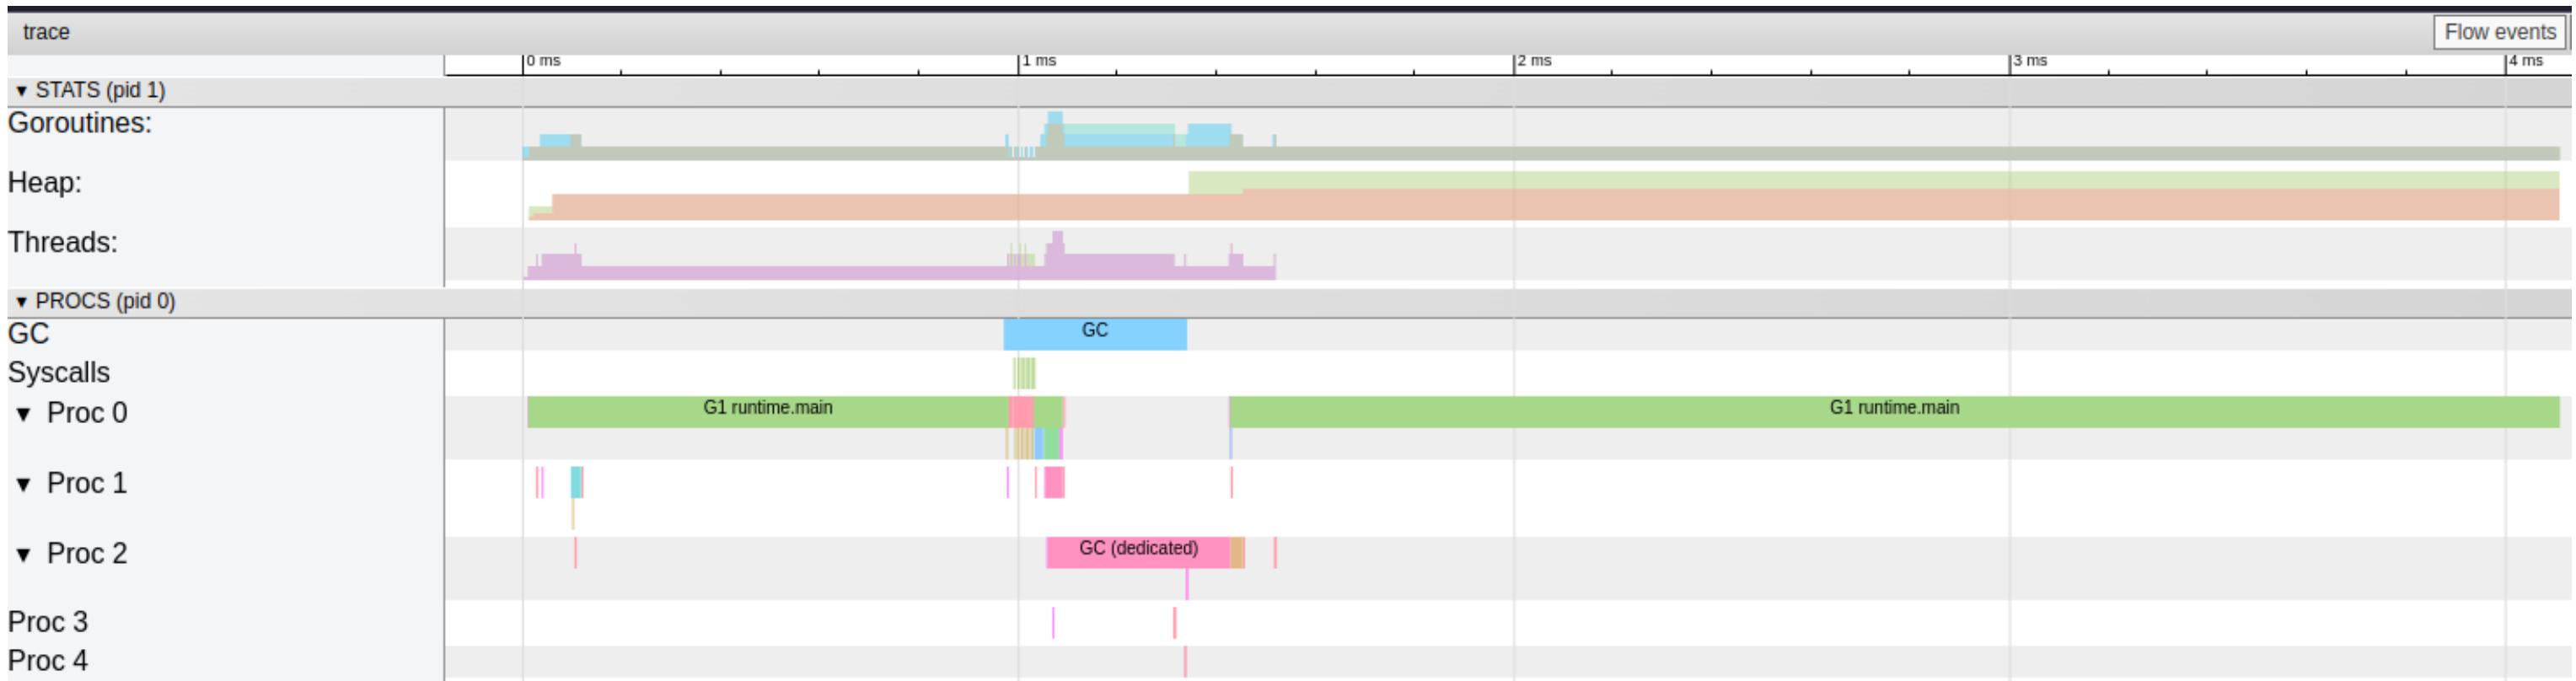
# Vanilla Code

```
slice := make([]CustomStruct, 0)
for i := 0; i < 100_000; i++ {
    slice = append(slice, CustomStruct{
        Name: "Alex",
        Age: 123,
    })
}
```

# Arena Code

```
mem := arena.NewArena()  
defer mem.Free()  
  
slice := arena.MakeSlice[CustomStruct](mem, 100_000, 100_000)  
  
for i := 0; i < 100_000; i++ {  
    c := arena.New[CustomStruct](mem)  
    c.Name = "Alex"  
    c.Age = 123  
    slice[i] = *c  
}
```

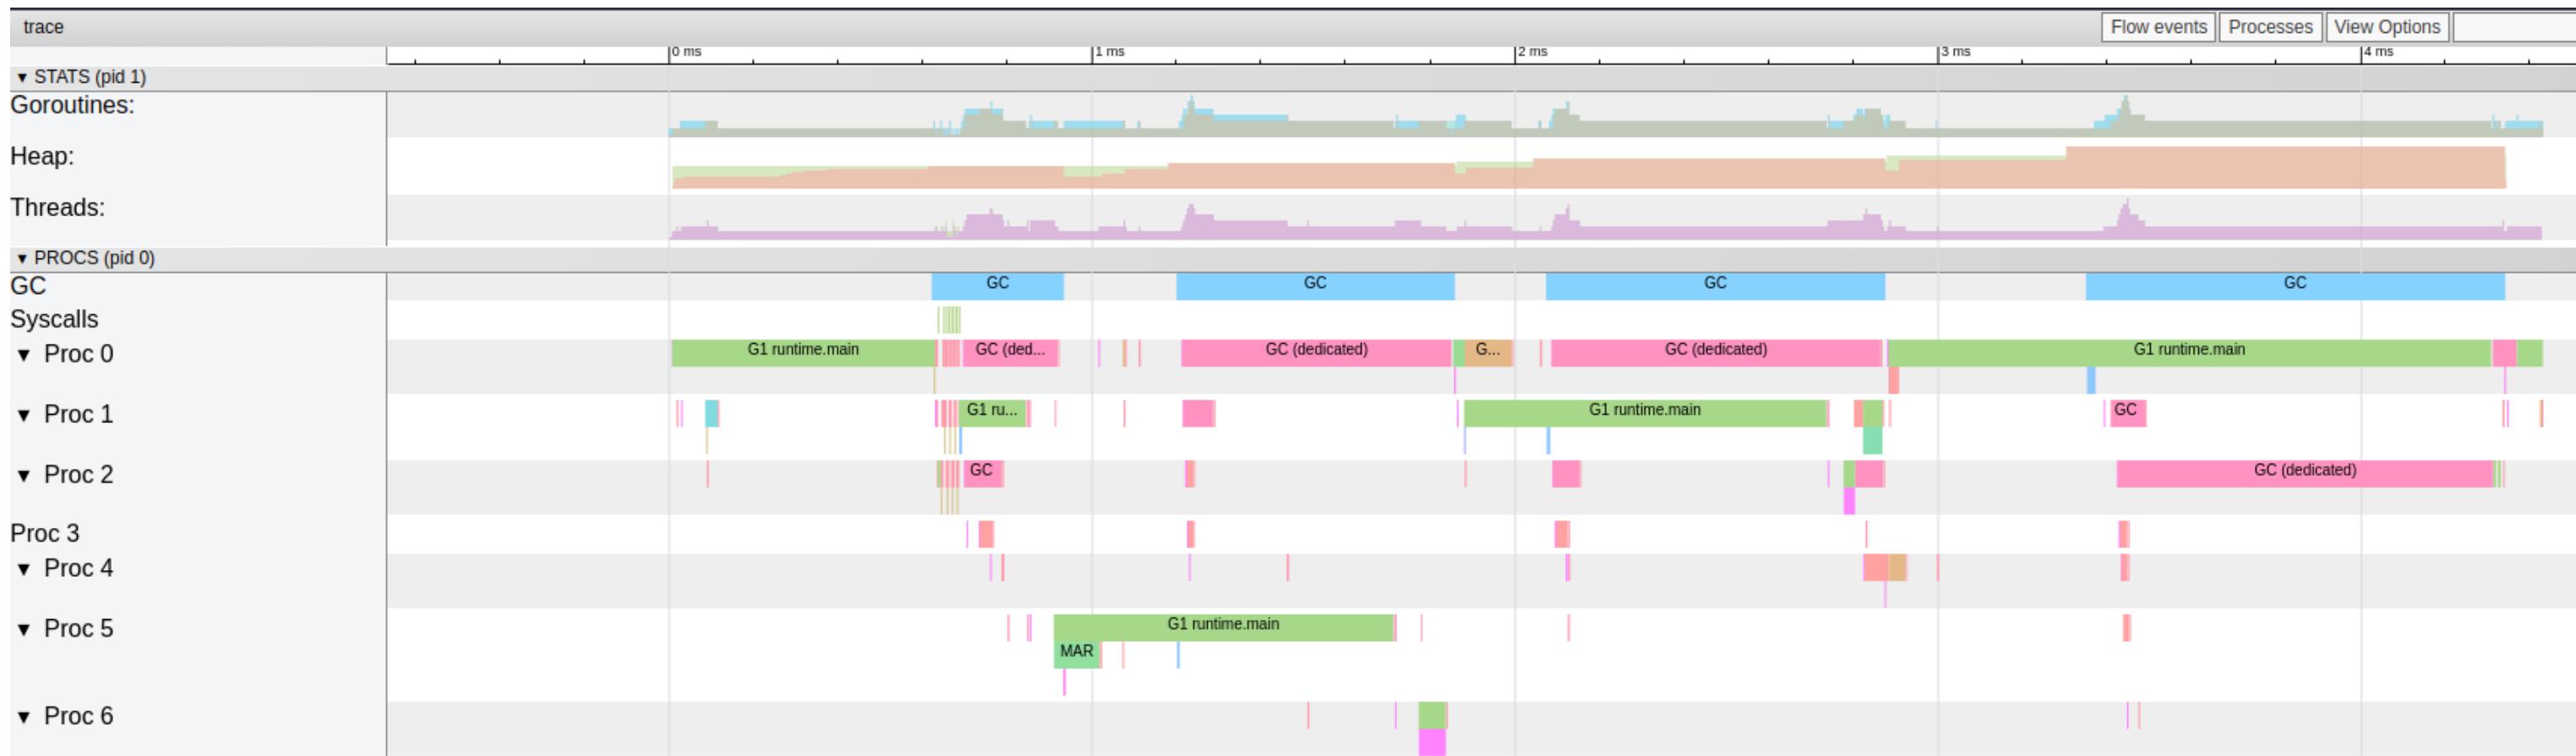
# Arena



# Arena

1 item selected.	Slice (1)						
Title	GC						
User Friendly Category	other						
Start	971,967 ns						
Wall Duration	369,701 ns						
Start Stack Trace	<table><thead><tr><th>Title</th></tr></thead><tbody><tr><td>runtime.newUserArenaChunk:713</td></tr><tr><td>runtime.(*userArena).refill:402</td></tr><tr><td>runtime.newUserArena:258</td></tr><tr><td>arena.runtime_arena_newArena:101</td></tr><tr><td>arena.NewArena:50</td></tr></tbody></table>	Title	runtime.newUserArenaChunk:713	runtime.(*userArena).refill:402	runtime.newUserArena:258	arena.runtime_arena_newArena:101	arena.NewArena:50
Title							
runtime.newUserArenaChunk:713							
runtime.(*userArena).refill:402							
runtime.newUserArena:258							
arena.runtime_arena_newArena:101							
arena.NewArena:50							

# Trace - Vanilla



# Trace - Vanilla

4 items selected. Slices (4)				
Name ▾	Wall Duration ▾	Self time ▾	Average Wall Duration ▾	Occurrences ▾
<u>GC</u>	2,763,703 ns	2,763,703 ns	690,926 ns	4
Selection start			622,416 ns	
Selection extent			3,719,647 ns	

# Benchmark

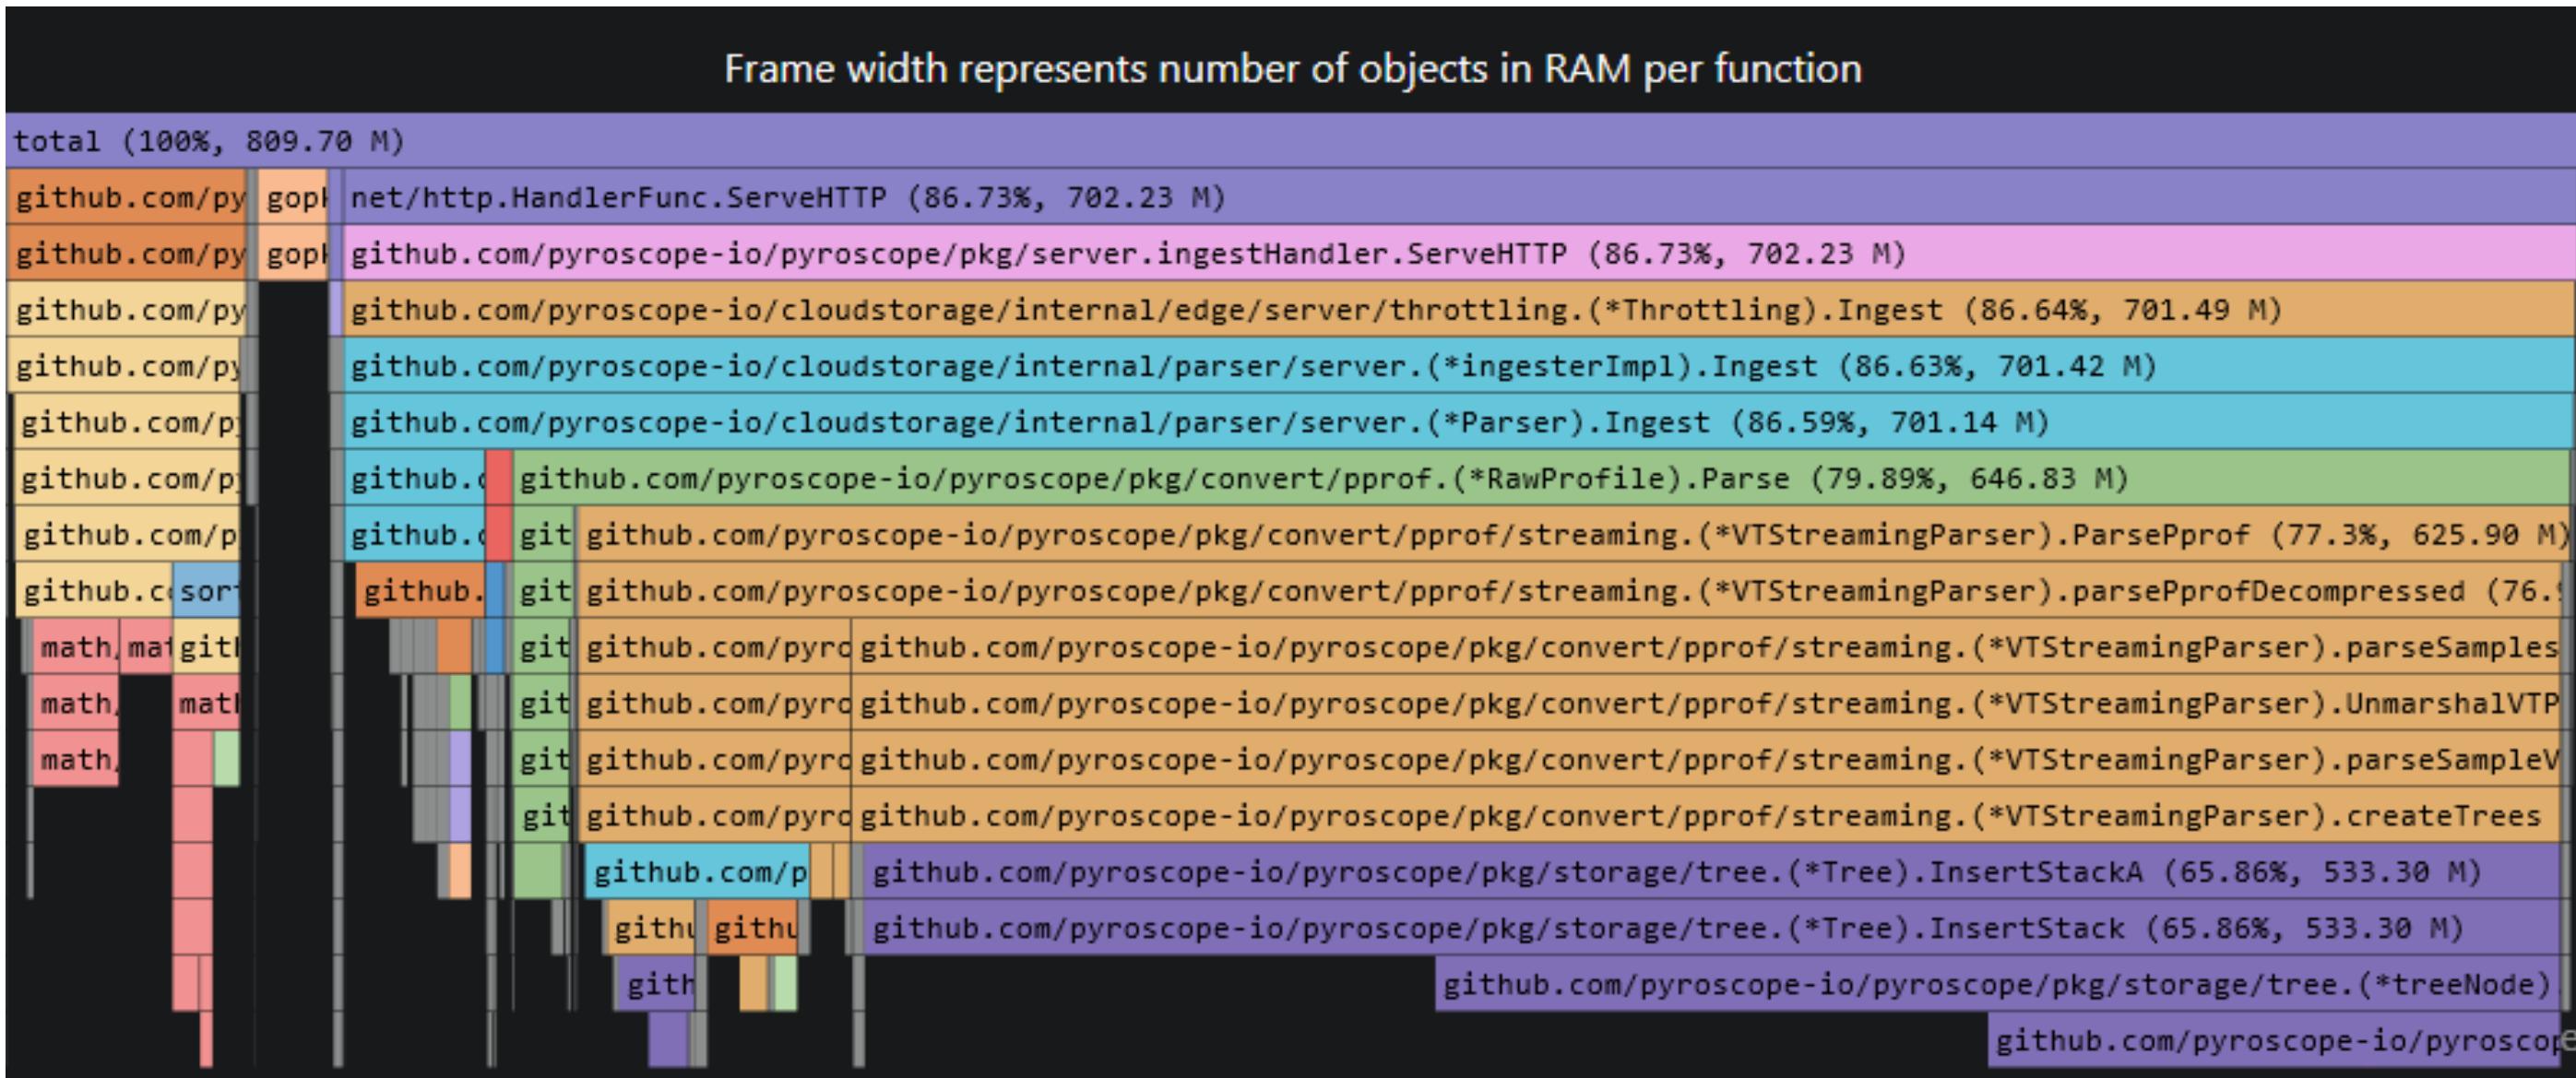
<b>BenchmarkVanilla-8</b>	<b>196</b>	<b>6240066 ns/op</b>	<b>14077976 B/op</b>	<b>29 allocs/op</b>
<b>BenchmarkArena-8</b>	<b>487</b>	<b>2433707 ns/op</b>	<b>4811826 B/op</b>	<b>3 allocs/op</b>

# Benchmark

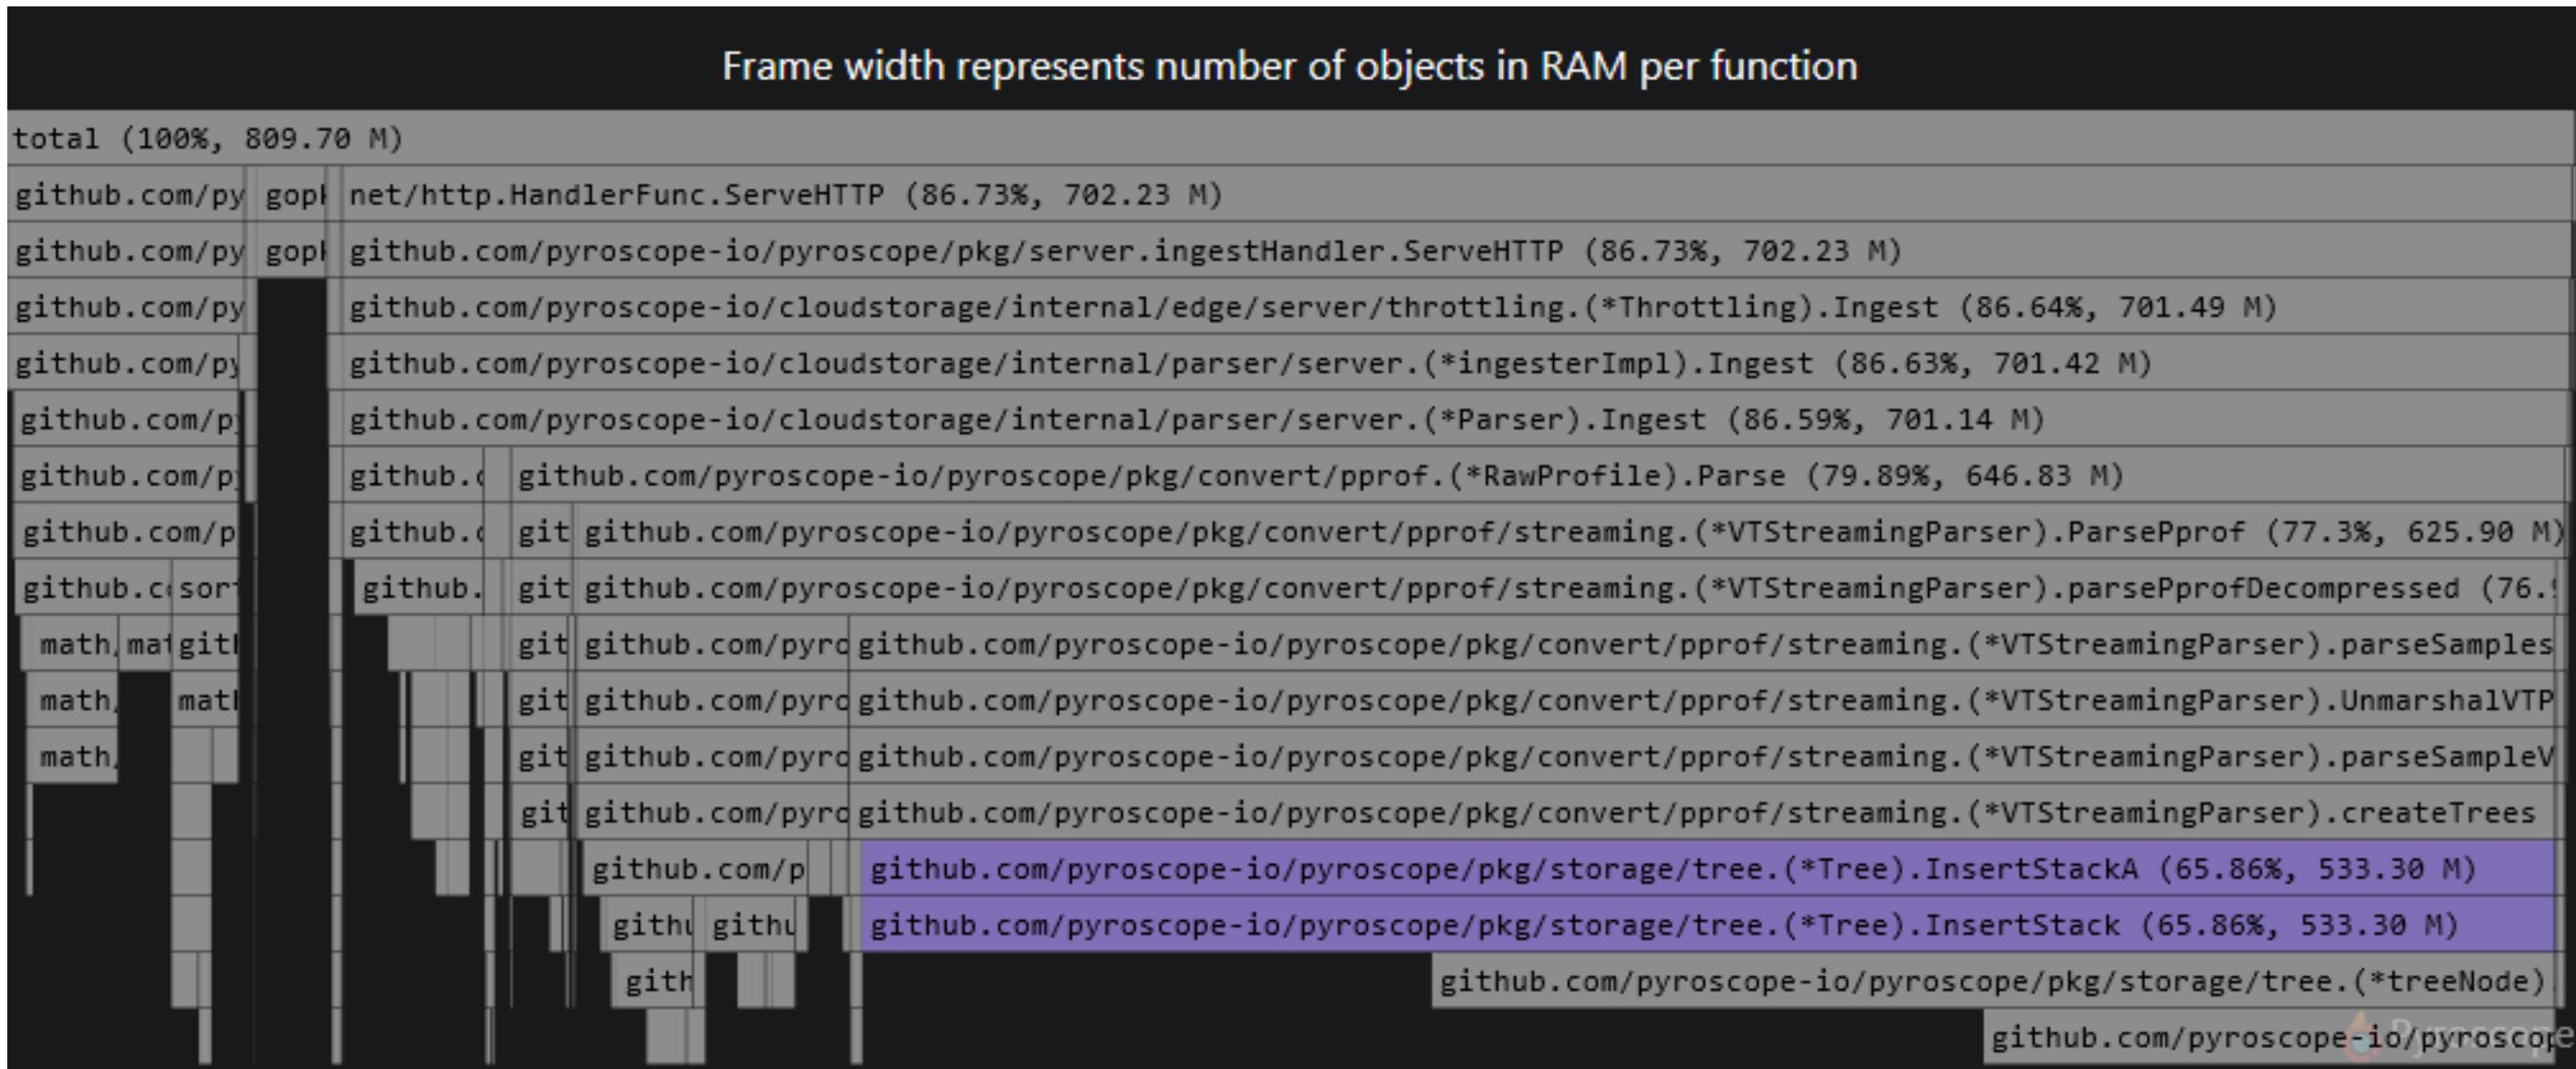
Nome	Iterações	ns/op	Byte/op	Alocações
Vanilla	196	6 240 066	14 077 976	29
Arena	487	2 433 707	4 811 826	3

# 4 - Case Study: Pyroscope

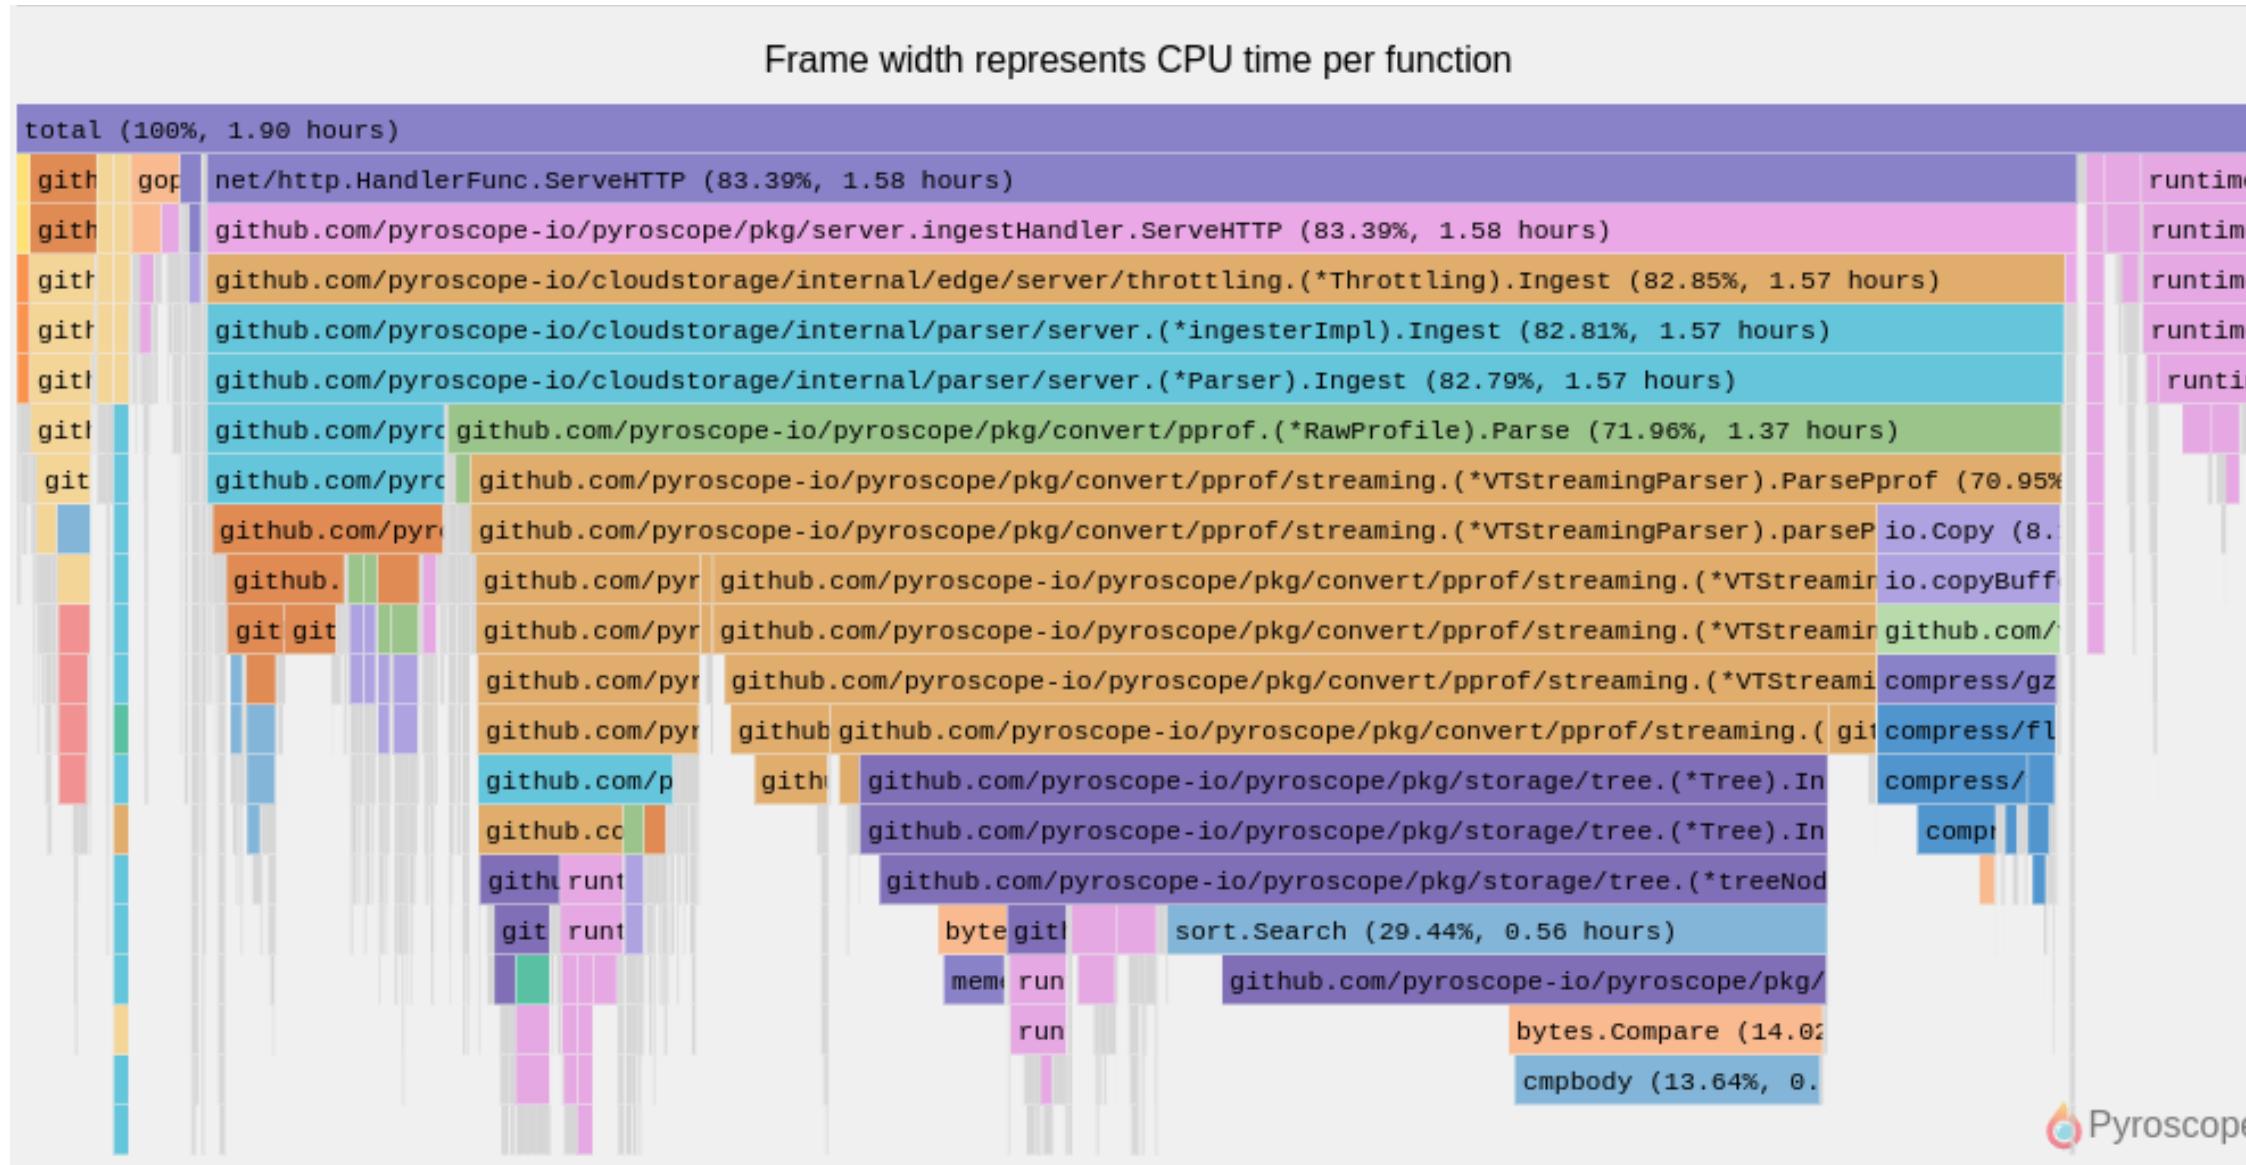
# Allocations profile



# Allocations profile

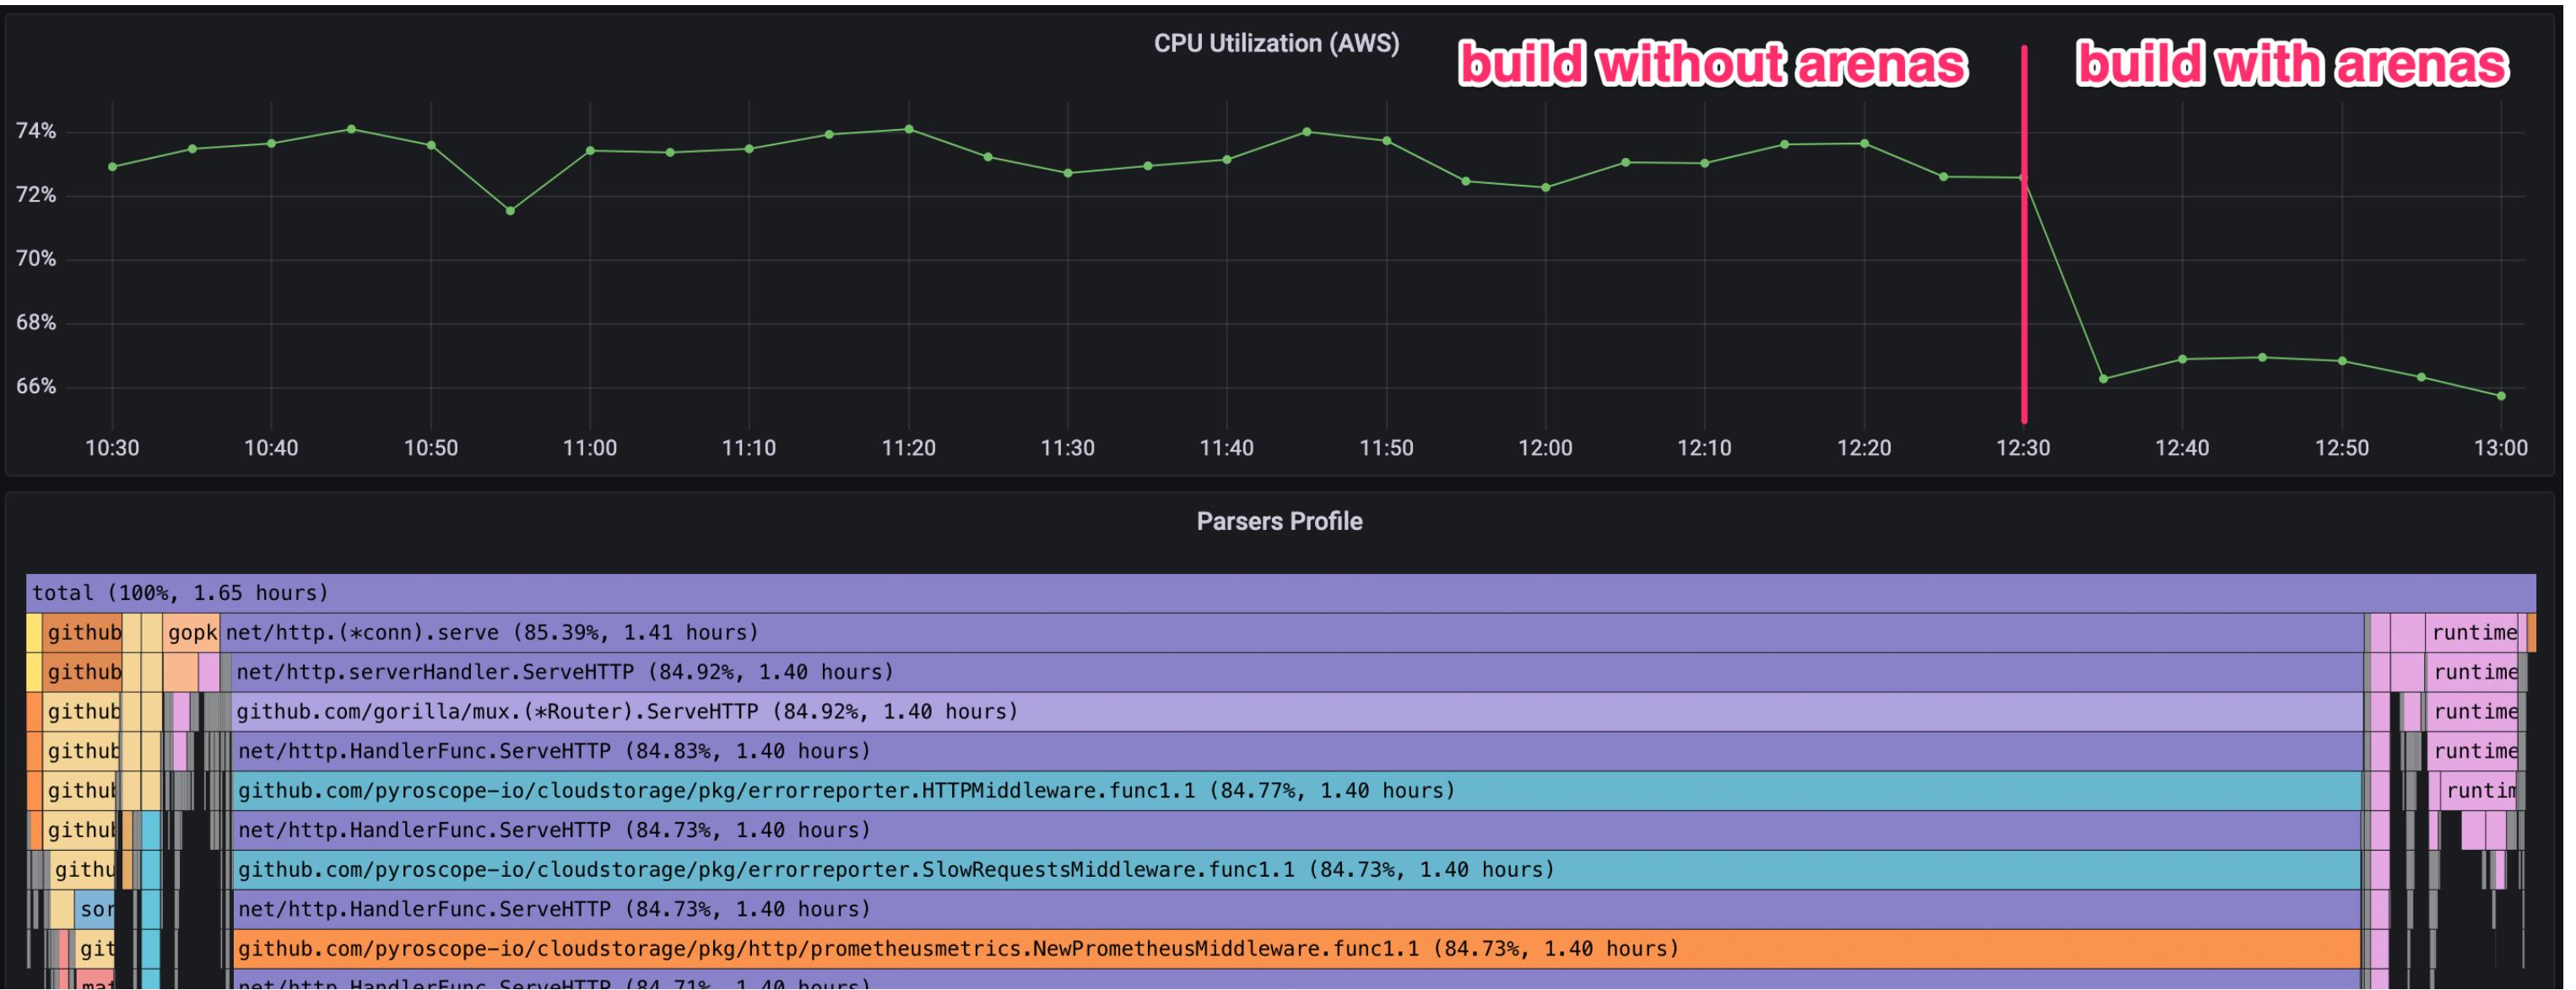


# CPU



```
total (100%, 1.90 hours)

net/http.HandlerFunc.ServeHTTP (83.39%, 1.58 hours)
github.com/pyroscope-io/pyroscope/pkg/server.ingestHandler.ServeHTTP (83.39%, 1.58 hours)
github.com/pyroscope-io/cloudstorage/internal/edge/server/throttling.(*Throttling).Ingest (82.85%, 1.57 hours)
github.com/pyroscope-io/cloudstorage/internal/parser/server.(*ingesterImpl).Ingest (82.81%, 1.57 hours)
github.com/pyroscope-io/cloudstorage/internal/parser/server.(*Parser).Ingest (82.79%, 1.57 hours)
github.com/pyroscope-io/pyroscope/pkg/convert/pprof.(*RawProfile).Parse (71.96%, 1.37 hours)
github.com/pyroscope-io/pyroscope/pkg/convert/pprof/streaming.(*VTStreamingParser).ParsePprof (70.95%, 1.35 hours)
github.com/pyroscope-io/pyroscope/pkg/convert/pprof/streaming.(*VTStreamingParser).parsePprofDecompressed (62.71%, 1.19 hours)
github.com/pyroscope-io/pyroscope/pkg/convert/pprof/streaming.(*VTStreamingParser).parseSamples (51.9%, 0.98 hours)
github.com/pyroscope-io/pyroscope/pkg/convert/pprof/streaming.(*VTStreamingParser).UnmarshalVTProfile (51.9%, 0.98 hours)
github.com/pyroscope-io/pyroscope/pkg/convert/pprof/streaming.(*VTStreamingParser).parseSampleVT (51.38%, 0.97 hours)
github.com/pyroscope-io/pyroscope/pkg/convert/pprof/streaming.(*VTStreamingParser).createTrees (44.49%, 0.84 hours)
github.com/pyroscope-io/pyroscope/pkg/storage/tree.(*Tree).InsertStackA (43.17%, 0.82 hours)
github.com/pyroscope-io/pyroscope/pkg/storage/tree.(*Tree).InsertStack (43.16%, 0.82 hours)
    github.com/pyroscope-io/pyroscope/pkg/storage/tree.(*treeNode).insert (42.28%, 0.80 hours)
        bytes.Equal (github.com/runtime/runtime.sort.Search (29.44%, 0.56 hours)
            memeqbody runtime. runti
                runtime.          github.com/pyroscope-io/pyroscope/pkg/storage/tree.(*treeNode).insert.func1 (26.97%, 0.51
                    bytes.Compare (14.02%, 0.27 hours)
                    cmpbody (13.64%, 0.26 hours)
```



# Oportunidades

# gRPC

Em gRPC existem muitos pequenos objetos sendo criados durante o encoding e decoding das mensagens.

<https://protobuf.dev/reference/cpp/arenas/>

# JSON decode/encode

Arena artesanal: [valyala](#) / [fastjson](#) (15x~ faster than stdlib)

# 5 - Guidelines

# Data driven?

Não? Você ta chutando.

# Several Allocations?



Não? Você provavelmente vai usar mais memória para fazer a mesma coisa.

# Same small structure?

yes? sync.Pool

# Hot Path? 🔥

Não? Otimização prematura.

# Vale a pena usar?

# Vale demais! porém...

- » Não é trivial.
- » Experimente combinações de GOGC e GOMEMLIMIT
- » Talvez você nunca precise usar. (e ta tudo bem)
- » Sua codebase que lida com a arena é muito delicada (leia-se possíveis crashes e memory leaks)



# Experimental API



- » Pode desaparecer
- » Ser descontinuada
- » Gerar retrabalho pelas mudanças de API

🙌 Obrigado! 🙌

[x.com/alextrending](https://x.com/alextrending)